

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.preprocessing import StandardScaler, MinMaxScaler

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, MaxPooling2D, Flatten

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: housing_data = pd.read_csv('BostonHousing.csv')
```

```
In [3]: housing_data.head()
```

```
Out[3]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

```
In [4]: housing_data.describe()
```

```
Out[4]:
```

	crim	zn	indus	chas	nox	rm	age
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000

```
In [5]: housing_data.isnull().sum()
```

```
Out[5]: crim      0
        zn        0
        indus     0
        chas      0
        nox       0
        rm        0
        age       0
        dis       0
        rad       0
        tax       0
        ptratio   0
        b         0
        lstat     0
        medv      0
        dtype: int64
```

```
In [6]: housing_data.duplicated().sum()
```

```
Out[6]: 0
```

```
In [7]: X = housing_data.drop(columns = ['medv'])
        y = housing_data.medv

        sc = StandardScaler()
        X = sc.fit_transform(X)

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [8]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[8]: ((354, 13), (152, 13), (354,), (152,))
```

```
In [9]: model = Sequential()

        model.add(Dense(128, input_shape=(13, ), activation='relu', name='dense_1'))
        model.add(Dense(64, activation='relu', name='dense_2'))
        model.add(Dense(1, activation='linear', name='dense_output'))

        model.compile(optimizer='adam', loss='mse', metrics=['mae'])
        model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 128)	1792
dense_2 (Dense)	(None, 64)	8256
dense_output (Dense)	(None, 1)	65
Total params: 10,113		
Trainable params: 10,113		
Non-trainable params: 0		

```
In [10]: history = model.fit(X_train, y_train, epochs=200, validation_split=0.2)
```

Epoch 1/200  
9/9 [=====] - 2s 39ms/step - loss: 580.0472 - mae: 22.2656 - val\_loss: 515.2385 - val\_mae: 20.8526  
Epoch 2/200  
9/9 [=====] - 0s 9ms/step - loss: 520.5151 - mae: 20.9143 - val\_loss: 448.8072 - val\_mae: 19.2797  
Epoch 3/200  
9/9 [=====] - 0s 10ms/step - loss: 447.1745 - mae: 19.1687 - val\_loss: 364.3033 - val\_mae: 17.1457  
Epoch 4/200  
9/9 [=====] - 0s 11ms/step - loss: 353.0789 - mae: 16.7185 - val\_loss: 263.5933 - val\_mae: 14.2895  
Epoch 5/200  
9/9 [=====] - 0s 9ms/step - loss: 245.7141 - mae: 13.4567 - val\_loss: 162.5475 - val\_mae: 10.7573  
Epoch 6/200  
9/9 [=====] - 0s 9ms/step - loss: 150.9113 - mae: 9.9895 - val\_loss: 88.9941 - val\_mae: 7.3551  
Epoch 7/200  
9/9 [=====] - 0s 9ms/step - loss: 85.1846 - mae: 7.4166 - val\_loss: 62.7974 - val\_mae: 5.8603  
Epoch 8/200  
9/9 [=====] - 0s 9ms/step - loss: 59.6940 - mae: 5.9908 - val\_loss: 56.2954 - val\_mae: 5.5551  
Epoch 9/200  
9/9 [=====] - 0s 9ms/step - loss: 47.3396 - mae: 5.2230 - val\_loss: 46.0090 - val\_mae: 4.9398  
Epoch 10/200  
9/9 [=====] - 0s 9ms/step - loss: 36.5004 - mae: 4.5176 - val\_loss: 35.0649 - val\_mae: 4.1577  
Epoch 11/200  
9/9 [=====] - 0s 9ms/step - loss: 28.9871 - mae: 4.0162 - val\_loss: 30.3657 - val\_mae: 3.8428  
Epoch 12/200  
9/9 [=====] - 0s 10ms/step - loss: 25.0969 - mae: 3.6780 - val\_loss: 27.6609 - val\_mae: 3.7041  
Epoch 13/200  
9/9 [=====] - 0s 10ms/step - loss: 22.5636 - mae: 3.4434 - val\_loss: 25.8646 - val\_mae: 3.6280  
Epoch 14/200  
9/9 [=====] - 0s 10ms/step - loss: 21.0843 - mae: 3.3112 - val\_loss: 24.1053 - val\_mae: 3.5487  
Epoch 15/200  
9/9 [=====] - 0s 10ms/step - loss: 19.5805 - mae: 3.1940 - val\_loss: 22.8075 - val\_mae: 3.4528  
Epoch 16/200  
9/9 [=====] - 0s 10ms/step - loss: 18.3879 - mae: 3.0638 - val\_loss: 21.6682 - val\_mae: 3.3447  
Epoch 17/200  
9/9 [=====] - 0s 9ms/step - loss: 17.5130 - mae: 2.9467 - val\_loss: 21.1239 - val\_mae: 3.2885  
Epoch 18/200  
9/9 [=====] - 0s 10ms/step - loss: 16.7063 - mae: 2.8636 - val\_loss: 20.5709 - val\_mae: 3.2310  
Epoch 19/200  
9/9 [=====] - 0s 9ms/step - loss: 16.0412 - mae: 2.7897 - val\_loss: 19.9443 - val\_mae: 3.1714  
Epoch 20/200  
9/9 [=====] - 0s 11ms/step - loss: 15.4640 - mae: 2.7447 - val\_loss: 19.0997 - val\_mae: 3.1049  
Epoch 21/200  
9/9 [=====] - 0s 9ms/step - loss: 14.9453 - mae: 2.6872 - val\_loss: 18.6990 - val\_mae: 3.0743  
Epoch 22/200



```
Epoch 193/200
9/9 [=====] - 0s 11ms/step - loss: 2.8564 - mae: 1.2454 -
val_loss: 7.4615 - val_mae: 2.0969
Epoch 194/200
9/9 [=====] - 0s 10ms/step - loss: 2.8677 - mae: 1.2607 -
val_loss: 7.2914 - val_mae: 2.0922
Epoch 195/200
9/9 [=====] - 0s 10ms/step - loss: 2.8928 - mae: 1.2356 -
val_loss: 7.2349 - val_mae: 2.0819
Epoch 196/200
9/9 [=====] - 0s 10ms/step - loss: 2.8208 - mae: 1.2274 -
val_loss: 7.3682 - val_mae: 2.0994
Epoch 197/200
9/9 [=====] - 0s 10ms/step - loss: 2.7770 - mae: 1.2369 -
val_loss: 7.3326 - val_mae: 2.0833
Epoch 198/200
9/9 [=====] - 0s 10ms/step - loss: 2.7763 - mae: 1.2236 -
val_loss: 7.2133 - val_mae: 2.0738
Epoch 199/200
9/9 [=====] - 0s 10ms/step - loss: 2.7261 - mae: 1.2082 -
val_loss: 7.2567 - val_mae: 2.0857
Epoch 200/200
9/9 [=====] - 0s 10ms/step - loss: 2.7398 - mae: 1.2067 -
val_loss: 7.2556 - val_mae: 2.0849
```

```
In [17]: print(len(history.history['mae']))

200
```

```
In [18]: sns.scatterplot(y = history.history['loss'],x = range(1,200+1))
sns.scatterplot(y = history.history['val_loss'],x = range(1,200+1))
plt.title('Training and Validation Loss')
plt.xlabel('epochs')
plt.ylabel('loss')
plt.legend(['Training Loss','Validation_Loss'])
plt.show()

sns.scatterplot(y = history.history['mae'],x = range(1,200+1))
sns.scatterplot(y = history.history['val_mae'],x = range(1,200+1))
plt.title('Training and Validation MAE')
plt.xlabel('epochs')
plt.ylabel('loss')
plt.legend(['Training Mae','Validation_Mae'])
plt.show()
```



```
In [19]: mse_nn, mae_nn = model.evaluate(X_test, y_test)
print('Mean squared error on test data: ', mse_nn)
print('Mean absolute error on test data: ', mae_nn)
```

```
5/5 [=====] - 0s 7ms/step - loss: 16.1675 - mae: 2.7948  
Mean squared error on test data: 16.16754150390625  
Mean absolute error on test data: 2.7947776317596436
```